

# Unsupervised learning for Consumer behaviour and marketing analysis

Stephanie Omwanda

## Introduction

Kira Plastinina is a Russian brand that is sold through a defunct chain of retail stores in Russia, Ukraine, Kazakhstan, Belarus, China, Philippines, and Armenia. The brand's Sales and Marketing team would like to understand their customer's behavior from data that they have collected over the past year. More specifically, they would like to learn the characteristics of customer groups.

## Research Question

Perform clustering and dimensionality reduction stating insights drawn from your analysis and visualizations. Upon implementation, provide comparisons between K-Means clustering vs Hierarchical clustering, highlighting the strengths and limitations of each approach in the context of your analysis. Your findings should help inform the team in formulating the marketing and sales strategies of the brand.

## Understanding the context

The dataset can be found here [ <http://bit.ly/EcommerceCustomersDataset> (<http://bit.ly/EcommerceCustomersDataset>) ]. The dataset consists of 10 numerical and 8 categorical attributes. The 'Revenue' attribute can be used as the class label. "Administrative", "Administrative Duration", "Informational", "Informational Duration", "Product Related" and "Product Related Duration" represents the number of different types of pages visited by the visitor in that session and total time spent in each of these page categories. The values of these features are derived from the URL information of the pages visited by the user and updated in real-time when a user takes an action, e.g. moving from one page to another. The "Bounce Rate", "Exit Rate" and "Page Value" features represent the metrics measured by "Google Analytics" for each page in the e-commerce site. The value of the "Bounce Rate" feature for a web page refers to the percentage of visitors who enter the site from that page and then leave ("bounce") without triggering any other requests to the analytics server during that session. The value of the "Exit Rate" feature for a specific web page is calculated as for all pageviews to the page, the percentage that was the last in the session. The "Page Value" feature represents the average value for a web page that a user visited before completing an e-commerce transaction. The "Special Day" feature indicates the closeness of the site visiting time to a specific special day (e.g. Mother's Day, Valentine's Day) in which the sessions are more likely to be finalized with the transaction. The value of this attribute is determined by considering the dynamics of e-commerce such as the duration between the order date and delivery date. For example, for Valentine's day, this value takes a nonzero value between February 2 and February 12, zero before and after this date unless it is close to another special day, and its maximum value of 1 on February 8. The dataset also includes the operating system, browser, region, traffic type, visitor type as returning or new visitor, a Boolean value indicating whether the date of the visit is weekend, and month of the year

## Metric of Success

The model's accuracy score is what will be used to measure the model's predictive power.

# Loading and Cleaning the dataset

Hide

```
# load libraries
library(readr) # provides a faster and friendly way to read rectangular data like
csvs
library(dplyr) # provides a flexible grammar of data manipulation library(tinytex)
library(knitr) # for dynamic report generation
options(warn = -1)
```

Hide

```
# Loading the csv file
df = read_csv('online_shoppers_intention.csv')
```

Parsed with column specification:

```
cols(
  Administrative = [32mcol_double()][39m,
  Administrative_Duration = [32mcol_double()][39m,
  Informational = [32mcol_double()][39m,
  Informational_Duration = [32mcol_double()][39m,
  ProductRelated = [32mcol_double()][39m,
  ProductRelated_Duration = [32mcol_double()][39m,
  BounceRates = [32mcol_double()][39m,
  ExitRates = [32mcol_double()][39m,
  PageValues = [32mcol_double()][39m,
  SpecialDay = [32mcol_double()][39m,
  Month = [31mcol_character()][39m,
  OperatingSystems = [32mcol_double()][39m,
  Browser = [32mcol_double()][39m,
  Region = [32mcol_double()][39m,
  TrafficType = [32mcol_double()][39m,
  VisitorType = [31mcol_character()][39m,
  Weekend = [33mcol_logical()][39m,
  Revenue = [33mcol_logical()][39m
)
```


Hide

```
# Previewing the first five rows of the dataframe
head(df)
```

Administrative <dbl>	Administrative_Duration <dbl>	Informational <dbl>	Informational_Duration <dbl>
0	0	0	0

0	0	0	0
0	-1	0	-1
0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

6 rows | 1-5 of 18 columns

[Hide](#)

```
# show structure of dataset  
str(df)
```

```
Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 12330 obs. of 18 variables:
 $ Administrative      : num  0 0 0 0 0 0 0 1 0 0 ...
 $ Administrative_Duration: num  0 0 -1 0 0 0 -1 -1 0 0 ...
 $ Informational       : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Informational_Duration: num  0 0 -1 0 0 0 -1 -1 0 0 ...
 $ ProductRelated      : num  1 2 1 2 10 19 1 1 2 3 ...
 $ ProductRelated_Duration: num  0 64 -1 2.67 627.5 ...
 $ BounceRates         : num  0.2 0 0.2 0.05 0.02 ...
 $ ExitRates           : num  0.2 0.1 0.2 0.14 0.05 ...
 $ PageValues          : num  0 0 0 0 0 0 0 0 0 0 ...
 $ SpecialDay          : num  0 0 0 0 0 0 0.4 0 0.8 0.4 ...
 $ Month               : chr   "Feb" "Feb" "Feb" "Feb" ...
 $ OperatingSystems    : num  1 2 4 3 3 2 2 1 2 2 ...
 $ Browser             : num  1 2 1 2 3 2 4 2 2 4 ...
 $ Region              : num  1 1 9 2 1 1 3 1 2 1 ...
 $ TrafficType         : num  1 2 3 4 4 3 3 5 3 2 ...
 $ VisitorType         : chr   "Returning_Visitor" "Returning_Visitor" "Returnin
g_Visitor" "Returning_Visitor" ...
 $ Weekend             : logi  FALSE FALSE FALSE FALSE TRUE FALSE ...
 $ Revenue             : logi  FALSE FALSE FALSE FALSE FALSE FALSE ...
- attr(*, "spec")=
 .. cols(
 ..   Administrative = [32mcol_double()][39m,
 ..   Administrative_Duration = [32mcol_double()][39m,
 ..   Informational = [32mcol_double()][39m,
 ..   Informational_Duration = [32mcol_double()][39m,
 ..   ProductRelated = [32mcol_double()][39m,
 ..   ProductRelated_Duration = [32mcol_double()][39m,
 ..   BounceRates = [32mcol_double()][39m,
 ..   ExitRates = [32mcol_double()][39m,
 ..   PageValues = [32mcol_double()][39m,
 ..   SpecialDay = [32mcol_double()][39m,
 ..   Month = [31mcol_character()][39m,
 ..   OperatingSystems = [32mcol_double()][39m,
 ..   Browser = [32mcol_double()][39m,
 ..   Region = [32mcol_double()][39m,
 ..   TrafficType = [32mcol_double()][39m,
 ..   VisitorType = [31mcol_character()][39m,
 ..   Weekend = [33mcol_logical()][39m,
 ..   Revenue = [33mcol_logical()][39m
 .. )
```

Hide

Hide

```
# catch a glimpse of the dataset
glimpse(df)
```

Observations: 12,330

Variables: 18

```
$ Administrative      [3m[38;5;246m<dbl>[39m[23m 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
$ Administrative_Duration [3m[38;5;246m<dbl>[39m[23m 0.0, 0.0, -1.0, 0.0, 0.0, 0.0
, -1.0, -1.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, ...
$ Informational        [3m[38;5;246m<dbl>[39m[23m 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
$ Informational_Duration [3m[38;5;246m<dbl>[39m[23m 0, 0, -1, 0, 0, 0, -1, -1, 0,
0, 0, 0, 0, 0, 0, 0, -1, 0, 0, 0, 0, -1, 0, 0...
$ ProductRelated       [3m[38;5;246m<dbl>[39m[23m 1, 2, 1, 2, 10, 19, 1, 1, 2,
3, 3, 16, 7, 6, 2, 23, 1, 13, 2, 20, 8, 1, 3, ...
$ ProductRelated_Duration [3m[38;5;246m<dbl>[39m[23m 0.000000, 64.000000, -1.00000
0, 2.666667, 627.500000, 154.216667, -1.000000...
$ BounceRates           [3m[38;5;246m<dbl>[39m[23m 0.2000000000, 0.0000000000, 0.2
000000000, 0.0500000000, 0.0200000000, 0.01578947...
$ ExitRates             [3m[38;5;246m<dbl>[39m[23m 0.2000000000, 0.1000000000, 0.2
000000000, 0.1400000000, 0.0500000000, 0.02456140...
$ PageValues            [3m[38;5;246m<dbl>[39m[23m 0.000000, 0.000000, 0.000000, 0.
000000, 0.000000, 0.000000, 0.000000, 0.0...
$ SpecialDay            [3m[38;5;246m<dbl>[39m[23m 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
0.4, 0.0, 0.8, 0.4, 0.0, 0.4, 0.0, 0.0, 0.0, ...
$ Month                 [3m[38;5;246m<chr>[39m[23m "Feb", "Feb", "Feb", "Feb", "
Feb", "Feb", "Feb", "Feb", "Feb", "Feb", "Feb"...
$ OperatingSystems      [3m[38;5;246m<dbl>[39m[23m 1, 2, 4, 3, 3, 2, 2, 1, 2, 2,
1, 1, 1, 2, 3, 1, 1, 1, 2, 2, 2, 3, 3, 2, 2, ...
$ Browser               [3m[38;5;246m<dbl>[39m[23m 1, 2, 1, 2, 3, 2, 4, 2, 2, 4,
1, 1, 1, 5, 2, 1, 1, 1, 2, 4, 2, 3, 2, 4, 2, ...
$ Region                [3m[38;5;246m<dbl>[39m[23m 1, 1, 9, 2, 1, 1, 3, 1, 2, 1,
3, 4, 1, 1, 3, 9, 4, 1, 1, 4, 5, 1, 1, 1, 4, ...
$ TrafficType           [3m[38;5;246m<dbl>[39m[23m 1, 2, 3, 4, 4, 3, 3, 5, 3, 2,
3, 3, 3, 3, 3, 3, 3, 4, 3, 4, 1, 3, 5, 3, 1, ...
$ VisitorType           [3m[38;5;246m<chr>[39m[23m "Returning_Visitor", "Returni
ng_Visitor", "Returning_Visitor", "Returning_V...
$ Weekend               [3m[38;5;246m<lgl>[39m[23m FALSE, FALSE, FALSE, FALSE, T
RUE, FALSE, FALSE, TRUE, FALSE, FALSE, FALSE, ...
$ Revenue               [3m[38;5;246m<lgl>[39m[23m FALSE, FALSE, FALSE, FALSE, F
ALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, ...
```

Hide

```
# checking for the statistical summary
summary(df)
```

Administrative	Administrative_Duration	Informational	Informational_Duration
ProductRelated			
Min. : 0.000	Min. : -1.00	Min. : 0.000	Min. : -1.00
Min. : 0.00			
1st Qu.: 0.000	1st Qu.: 0.00	1st Qu.: 0.000	1st Qu.: 0.00
1st Qu.: 7.00			
Median : 1.000	Median : 8.00	Median : 0.000	Median : 0.00
Median : 18.00			
Mean : 2.318	Mean : 80.91	Mean : 0.504	Mean : 34.51
Mean : 31.76			
3rd Qu.: 4.000	3rd Qu.: 93.50	3rd Qu.: 0.000	3rd Qu.: 0.00
3rd Qu.: 38.00			
Max. : 27.000	Max. : 3398.75	Max. : 24.000	Max. : 2549.38
Max. : 705.00			
NA's : 14	NA's : 14	NA's : 14	NA's : 14
NA's : 14			
ProductRelated_Duration	BounceRates	ExitRates	PageValues
SpecialDay			
Min. : -1.0	Min. : 0.000000	Min. : 0.000000	Min. : 0.000
n. : 0.000000			
1st Qu.: 185.0	1st Qu.: 0.000000	1st Qu.: 0.01429	1st Qu.: 0.000
t Qu.: 0.000000			
Median : 599.8	Median : 0.003119	Median : 0.02512	Median : 0.000
dian : 0.000000			
Mean : 1196.0	Mean : 0.022152	Mean : 0.04300	Mean : 5.889
an : 0.06143			
3rd Qu.: 1466.5	3rd Qu.: 0.016684	3rd Qu.: 0.05000	3rd Qu.: 0.000
d Qu.: 0.000000			
Max. : 63973.5	Max. : 0.200000	Max. : 0.20000	Max. : 361.764
x. : 1.00000			
NA's : 14	NA's : 14	NA's : 14	
Month	OperatingSystems	Browser	Region
VisitorType			
Length:12330	Min. : 1.000	Min. : 1.000	Min. : 1.000
0 Length:12330			
Class :character	1st Qu.: 2.000	1st Qu.: 2.000	1st Qu.: 1.000
0 Class :character			
Mode :character	Median : 2.000	Median : 2.000	Median : 3.000
0 Mode :character			
	Mean : 2.124	Mean : 2.357	Mean : 3.147
7			
	3rd Qu.: 3.000	3rd Qu.: 2.000	3rd Qu.: 4.000
0			
	Max. : 8.000	Max. : 13.000	Max. : 9.000
0			
Weekend	Revenue		
Mode :logical	Mode :logical		
FALSE:9462	FALSE:10422		
TRUE :2868	TRUE :1908		

Hide

```
# determine the dimensions of the dataset
dim(df)
```

```
[1] 12330    18
```

From the chunk above, it is evident that the dataset contains 12,330 observations of 18 variables

Hide

```
# checking if there exists null values by calculating the sum of the null values per column
colSums((is.na(df)))
```


Administrative_Duration	Administrative_Duration	Informational_Duration	Informational_Duration
14	14	14	14
ProductRelated_ExitRates	ProductRelated_Duration	BounceRates	
14	14	14	
PageValues	SpecialDay	Month	OpenPagesPerSession
0	0	0	0
Browser	Region	TrafficType	
0	0	0	
Weekend	Revenue		
0	0		

We can see that about 8 variables have the same number of missing values

Hide



```
# a plot showing missing values
library(DataExplorer) # simplifies and automates EDA processes and aids in report
generation
plot_missing(df)
```


[Hide](#)

```
# Dropping missing values
df = na.omit(df)
```

[Hide](#)

```
colSums((is.na(df)))
```

Administrative	Administrative_Duration	Informational	Informational_Duration
0	0	0	0
ProductRelated	ProductRelated_Duration	BounceRates	
0	0	0	
PageValues	SpecialDay	Month	OpenPagesPerVisit
0	0	0	
Browser	Region	TrafficType	
0	0	0	
Weekend	Revenue		
0	0		

The missing values have been successfully omitted

Hide

```
# Checking for duplicated data
anyDuplicated(df)
```

```
[1] 159
```

The dataset is seen to contain a number of duplicates which will mess up with analysis and the prediction model, they will therefore be dealt with in the next chunk

Hide

```
# Dropping duplicates
df = distinct(df)

# confirming whether the drop was successful
anyDuplicated(df)
```

```
[1] 0
```

Our dataset is now free of redundant data.

Hide

```
# Checking the type of the dataset
class(df)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

Hide

```
# Changing the type of the loaded dataset to a dataframe
df = as.data.frame(df)
class(df)
```

```
[1] "data.frame"
```

Hide

```
library(magrittr) # offers a set of operators that provide semantics that will improve code
# Checking the datatypes for each column
columns = colnames(df)
for (column in seq(length(colnames(df)))){
  print(columns[column])
  print(class(df[, column]))
  cat('\n')
}
```

```
[1] "Administrative"
[1] "numeric"

[1] "Administrative_Duration"
[1] "numeric"

[1] "Informational"
[1] "numeric"

[1] "Informational_Duration"
[1] "numeric"

[1] "ProductRelated"
[1] "numeric"

[1] "ProductRelated_Duration"
[1] "numeric"
```

```
[1] "BounceRates"
[1] "numeric"

[1] "ExitRates"
[1] "numeric"

[1] "PageValues"
[1] "numeric"

[1] "SpecialDay"
[1] "numeric"

[1] "Month"
[1] "character"

[1] "OperatingSystems"
[1] "numeric"

[1] "Browser"
[1] "numeric"

[1] "Region"
[1] "numeric"

[1] "TrafficType"
[1] "numeric"

[1] "VisitorType"
[1] "character"

[1] "Weekend"
[1] "logical"

[1] "Revenue"
[1] "logical"
```

[Hide](#)


```
# Cleaning column names, by making them uniform
colnames(df) = tolower(colnames(df))
```

Now that our dataset is free of missing values and duplicates, we will go ahead and do a visual check on outliers in the dataset

[Hide](#)


```
library(ggplot2) # aids in creation of data visuals with the help of grammar of graphics
# Plotting boxplots
options(repr.plot.width = 11, repr.plot.height = 5)
ggplot(df, aes(month, productrelated, col = weekend)) +
  geom_boxplot() +
  labs(x = 'Month', y = 'Product related', title = 'Outliers : product related and the month') +
  theme(legend.position = 'top', legend.text = element_text(size = 10),
        plot.title = element_text(size = 14, color = 'black', face = 'bold'))
```

## Outliers : product related and the month


[Hide](#)

```
# Plotting boxplots to check for outliers
options(repr.plot.width = 7, repr.plot.height = 5)
ggplot(df, aes(visitor type, productrelated, col = revenue)) +
  geom_boxplot() +
  labs(x = 'Visitor type', y = 'Product related', title = 'Outliers : Product related and the visitor type') +
  scale_color_brewer(palette = 'Set1') +
  theme(legend.position = 'top',
        plot.title = element_text(size = 14, color = 'black', face = 'bold'))
```

## Outliers : Product related and the visitor type



## Exploratory data analysis

This is where we explore the data so as to: \* maximize insights on the data set \* uncover underlying structure \* extract important variables \* test underlying assumptions \* develop models with great explanatory predictive power \* determine optimal factor settings

Here we will perform : \* univariate analysis \* bivariate analysis \* multivariate analysis

In EDA we have both graphical and non-graphical analysis, where non graphical contains Measures of central tendencies :(Mean, mode and median for numerical data and Mode for categorical data) and Measures of dispersion. Because the summary function did most of the non graphical analysis, this section will be densely populated with visualizations and their analyses

# Univariate analysis

[Hide](#)

```
library(gridExtra) # provides a number of user level functions to work with grid graphics

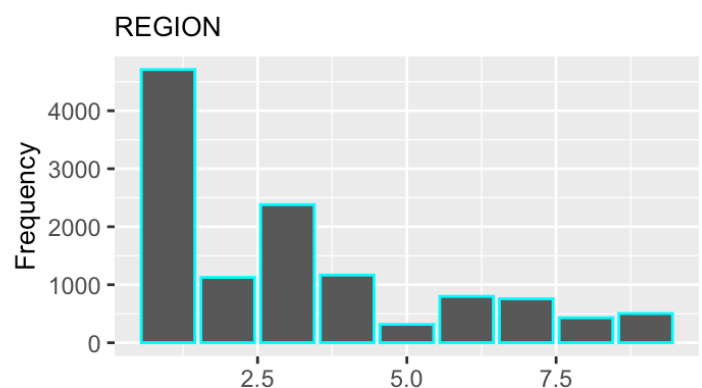

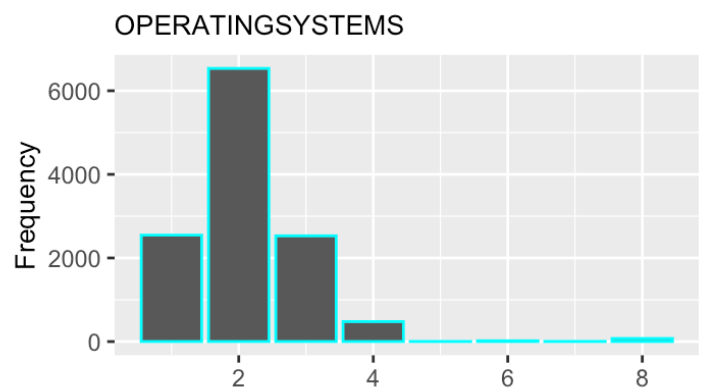

# Plotting bar plots showing frequency for each of variables
# we will plot the first 4 in this current chunk then the next for on the next chunk to avoid crampling up and untidyness

fac_cols = c('month', 'operatingsystems', 'browser', 'region')

columns = colnames(select(df, fac_cols))

p = list()
options(repr.plot.width = 10, repr.plot.height = 6)
for (i in 1:4){
  p[[i]] = ggplot(df, aes_string(columns[i])) + geom_bar(color = 'cyan') + labs(y = 'Frequency', x = '', title = toupper(columns[i])) +
    theme(plot.title = element_text(size = 10),
          axis.title.y = element_text(size = 10))
}

do.call(grid.arrange, p)
```


[Hide](#)

```




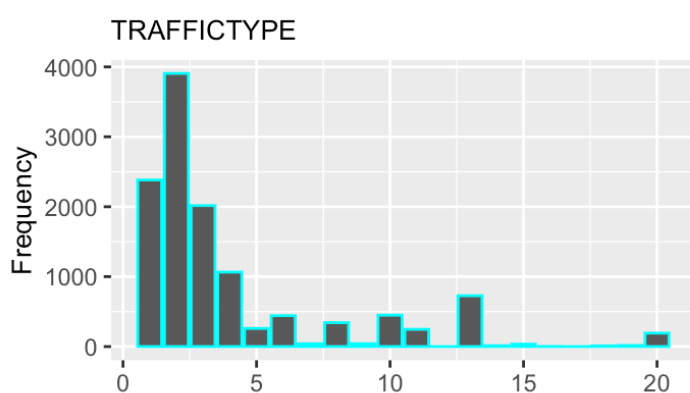
fac_cols_2 = c('traffictype', 'visitortype', 'weekend', 'revenue')

columns_2 = colnames(select(df, fac_cols_2))

p_2 = list()
options(repr.plot.width = 10, repr.plot.height = 6)
for (i in 1:4){
  p_2[[i]] = ggplot(df, aes_string(columns_2[i])) + geom_bar(color = 'cyan') + lab
s(y = 'Frequency', x = '', title = toupper(columns_2[i])) +
  theme(plot.title = element_text(size = 10),
        axis.title.y = element_text(size = 10))
}

do.call(grid.arrange, p_2)

```



Now that we have seen the frequency of each column in the dataset, we will check for their distribution and determine whether they are negatively or positively skewed

Hide



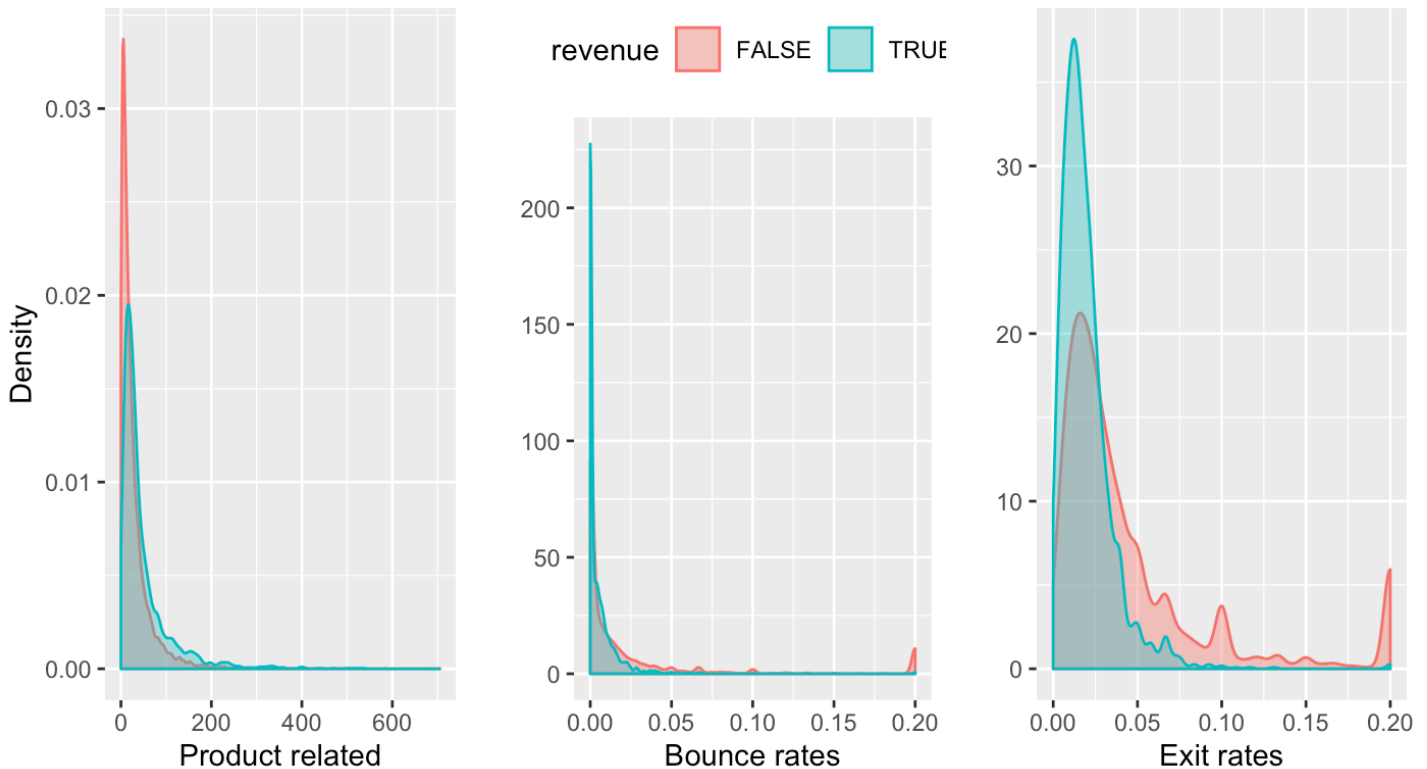
```
library(grid)
# Plotting density plots to check for distributions
options(repr.plot.width = 11, repr.plot.height = 5)
p1 = ggplot(df, aes(productrelated, col = revenue)) +
  geom_density(aes(fill = revenue), alpha = 0.4) +
  labs(x = 'Product related', y = 'Density', title = '') +
  theme(legend.position = 'none',
        plot.title = element_text(size = 12))

p2 = ggplot(df, aes(bouncerates, col = revenue)) +
  geom_density(aes(fill = revenue), alpha = 0.4) +
  labs(x = 'Bounce rates', y = '', title = '') +
  theme(legend.position = 'top')

p3 = ggplot(df, aes(exitrates, col = revenue)) +
  geom_density(aes(fill = revenue), alpha = 0.4) +
  labs(x = 'Exit rates', y = '', title = '') +
  theme(legend.position = 'none',
        plot.title = element_text(size = 12))

grid.arrange(p1, p2, p3, ncol = 3, top = textGrob("Density plots showing distribut
ion",gp=gpar(fontsize=13,font=3, color = 'black')))
```

*Density plots showing distribution*



So far, the variables we are working with are all positively skewed

# Bivariate analysis

[Hide](#)

```
# Plotting scatter plots to check for correlations
options(repr.plot.width = 11, repr.plot.height = 5)

p1 = ggplot(df, aes(productrelated, productrelated_duration, col = revenue)) +
  geom_point() + theme(legend.position = 'none') +
  labs(x='Product related', y='Product related duration')

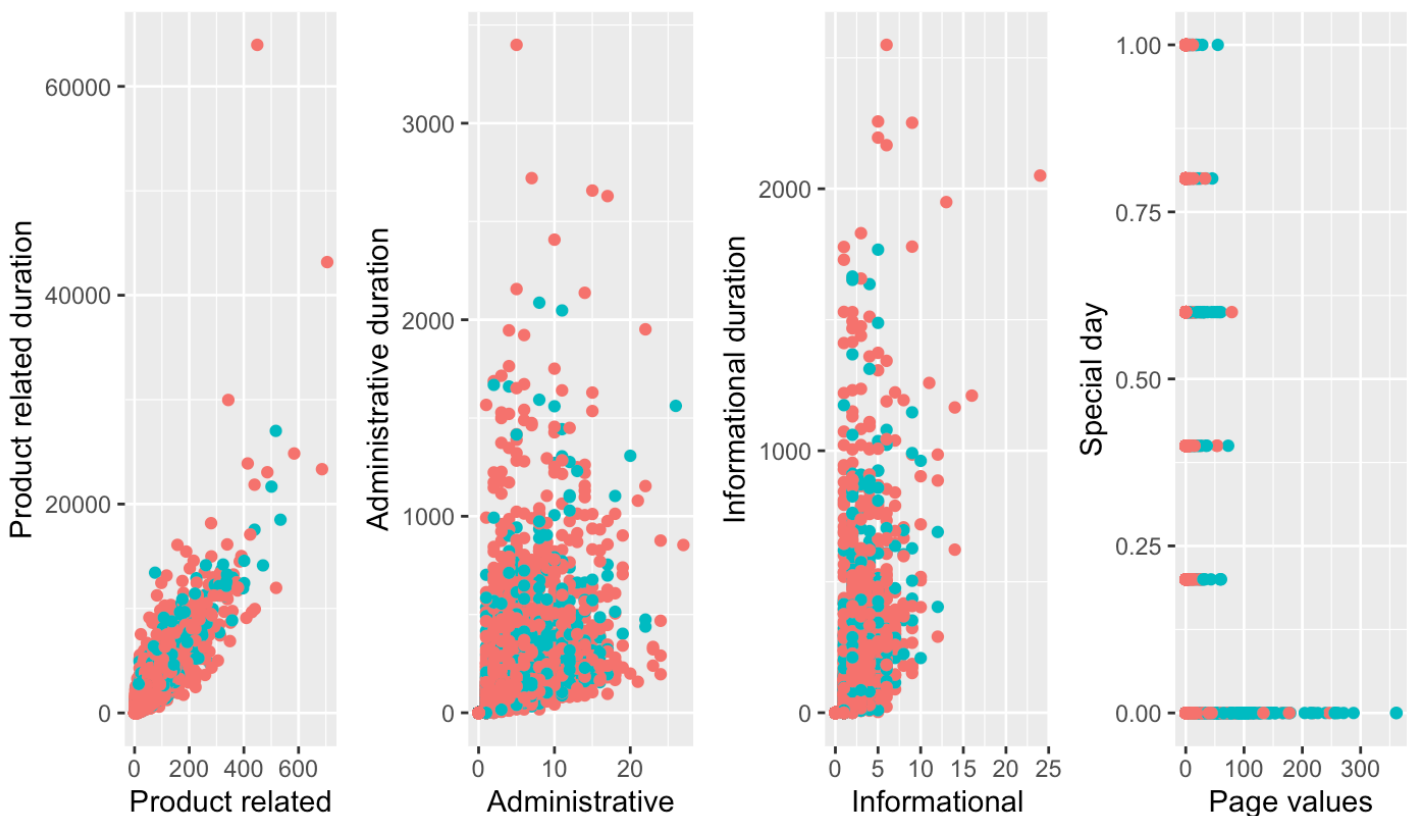
p2 = ggplot(df, aes(administrative, administrative_duration, col = revenue)) +
  geom_point() + theme(legend.position = 'none') +
  labs(x = 'Administrative', y = 'Administrative duration')

p3 = ggplot(df, aes(informational, informational_duration, col = revenue)) +
  geom_point() + theme(legend.position = 'none') +
  labs(x = 'Informational', y = 'Informational duration')

p4 = ggplot(df, aes(pagevalues, specialday , col = revenue)) +
  geom_point() + theme(legend.position = 'none') +
  labs(x = 'Page values', y = 'Special day')

grid.arrange(p1, p2, p3, p4, ncol = 4,
              top = textGrob("Scatter plots to show correlations",gp=gpar(fontsize=
14,font=3, color = 'darkmagenta')))
```

*Scatter plots to show correlations*



[Hide](#)

```
library (ggExtra) # used to add marginal histograms, densityplots and boxplots scatter plots

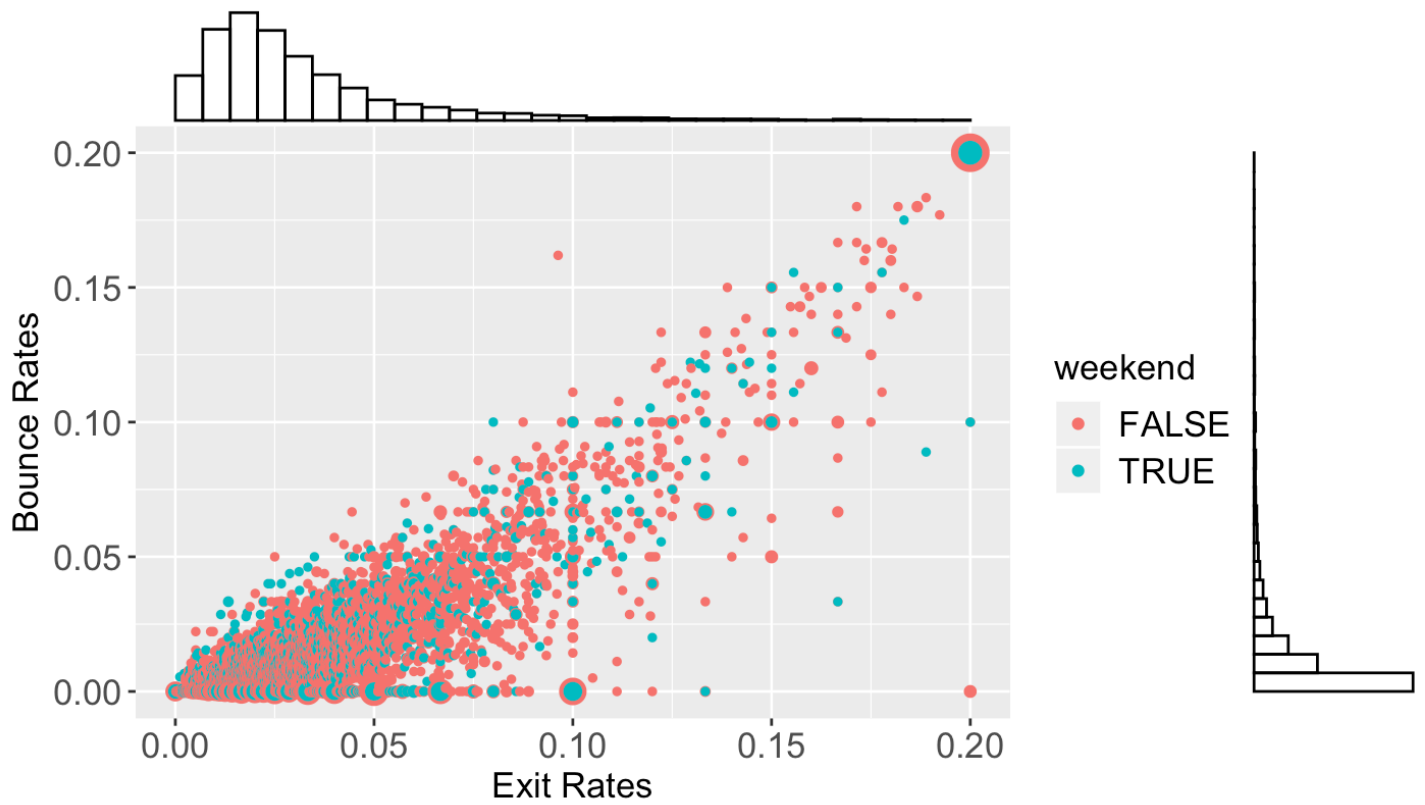
# Plotting scatterplot with marginal density plots (default) or histograms using ggMarginal

options(repr.plot.width = 7, repr.plot.height = 5)

g = ggplot(data =df, aes(x =exitrates, y = bouncerates, col = weekend)) +
  geom_count(show.legend=c(size=FALSE)) +
  labs(title = 'Bounce Rates Vs Exit Rates', y = 'Bounce Rates', x = 'Exit Rates') +
  theme(plot.title = element_text(size = 14, face = 'bold'),
        axis.title.x = element_text(size = 13),
        axis.title.y = element_text(size = 13),
        axis.text.x = element_text(size = 13),
        axis.text.y = element_text(size = 13),
        legend.title = element_text(size = 13),
        legend.text = element_text(size = 13))


ggMarginal(g, type = "histogram", fill="transparent")
```

## Bounce Rates Vs Exit Rates

[Hide](#)


```
ggMarginal(g, type = "boxplot", fill="transparent")
```

## Bounce Rates Vs Exit Rates

[Hide](#)

```
ggMarginal(g, type = "density", fill="transparent")
```

## Bounce Rates Vs Exit Rates




## Multivariate analysis

[Hide](#)

```
library(ggcorrplot)
# Plotting a correlogram to check for correlations among all the features
options(repr.plot.width = 6, repr.plot.height = 5)

corr = round(cor(select_if(df, is.numeric)), 2)
ggcorrplot(corr, hc.order = T, ggtheme = ggplot2::theme_gray,
  colors = c("cyan", "peachpuff4", "pink"), lab = F)
```



Now that we have created visual representations of how the dataset looks like we will go ahead and perform analyses and create the models we need to make the best predictions

## K-Means

Because K-means is a type unsupervised-learning, the class attribute is not needed for the execution of the algorithm. In the 'Understanding the Context' section, we mentioned that the 'Revenue' attribute will be used as the class label. It will be removed and stored in another variable. Afterwards we can normalize the attributes left.

Hide

```
# creating a new dataframe
df.new <- df[,c(1:17)]
colnames(df.new)
```

```
[1] "administrative"      "administrative_duration" "informational"
"informational_duration"
[5] "productrelated"      "productrelated_duration" "bouncerrates"
"exitrates"
[9] "pagevalues"          "specialday"              "month"
"operatingsystems"
[13] "browser"             "region"                  "traffictype"
"visitortype"
[17] "weekend"
```

Hide

```
# storing the class attribute in another variable
df.class <- df["revenue"]
colnames(df.class)
```

```
[1] "revenue"
```

Hide

```
# Normalising the data so that no particular attribute has more impact on clusteri
ng than others
normalize <- function(x){
  return ((x-min(x)) / (max(x)-min(x)))
}
df.new$administrative <- normalize(df.new$administrative)
df.new$administrative_duration <- normalize(df.new$administrative_duration)
df.new$infromational <- normalize(df.new$informational)
df.new$infromational_duration <- normalize(df.new$informational_duration)
df.new$productrelated <- normalize(df.new$productrelated)
df.new$productrelated_duration <- normalize(df.new$productrelated_duration)
df.new$bouncerrates<- normalize(df.new$bouncerrates)
df.new$exitrates<- normalize(df.new$exitrates)
df.new$pagevalues<- normalize(df.new$pagevalues)
df.new$specialday<- normalize(df.new$specialday)
df.new$opertaingsystem<- normalize(df.new$operatingsystems)
df.new$browser<- normalize(df.new$browser)
df.new$region<- normalize(df.new$region)
df.new$traffictype<- normalize(df.new$traffictype)

head(df)
```

	<b>administrative</b> <dbl>	<b>administrative_duration</b> <dbl>	<b>informational</b> <dbl>	<b>informational_duration</b> <dbl>
1	0	0	0	0
2	0	0	0	0



3	0	-1	0	-1
4	0	0	0	0
5	0	0	0	0
6	0	0	0	0

6 rows | 1-6 of 18 columns

Hide

NA

Hide

```
# Storing numeric values in a new variable
df.clust <- df.new [,c("administrative","administrative_duration","informational",
"informational_duration","productrelated","productrelated_duration","bouncerrates",
"exitrates","pagevalues","specialday","operatingsystems","browser","region","traff
ictype")]
```

```
# Apply k means clustering algorithm with a number of centroids k
result<- kmeans(df.clust,3)
```

Hide

```
# preview the number of records in each cluster

result$size
```

```
[1] 116 11530 553
```

Hide

```
# get the value of cluster center datapoint value(k centers for k)
result$centers
```

```

administrative administrative_duration informational informational_duration prod
uctrelated
1      0.22701149      0.08298299      4.4655172      1156.40264
0.14458303
2      0.07877357      0.02160253      0.3204683      8.72358
0.04184225
3      0.22182037      0.06867505      3.6057866      344.04175
0.10038347
productrelated_duration bounce rates exit rates page values special day operating sys
tems      browser      region
1      0.08241265 0.04112800 0.1182392 0.01846520 0.03793103      2.12
0690 0.10991379 0.2198276
2      0.01701258 0.10601490 0.2132804 0.01603853 0.06374675      2.12
5412 0.11389853 0.2709128
3      0.04471927 0.03621331 0.1053481 0.02469705 0.03001808      2.10
3074 0.09885473 0.2429928
traffic type
1      0.1279492
2      0.1632081
3      0.1400019

```

[Hide](#)

```


# getting cluster vector that shows the cluster where each record falls
result$cluster

```


Hide

Hide


```
# Plot dist in clusters  
plot(df.clust[c(1,2)], col = result$cluster)
```

[Hide](#)

```
# Plot dist in clusters  
plot(df.clust[c(3,4)], col = result$cluster)
```

[Hide](#)

```
# Plot dist in clusters  
plot(df.clust[c(5,6)], col = result$cluster)
```

[Hide](#)

```
NA
NA
```

[Hide](#)

```
df.class[] <- lapply(df.class, as.numeric)
sapply(df.class, class)
```

```
revenue
"numeric"
```

[Hide](#)

```
#install.packages('useful')
#library(useful)

#df.clust[] <- lapply(df.clust, as.numeric)
#plot(df.clust[c(1,2)], col = df.class)
```

[Hide](#)

```
# result table
#table(result$cluster, df.class)
```

Hide

```
# accuracy score
mean(df.new == result$cluster)
```

```
[1] 0.02897779
```

## Hierachical Clustering

Hide

```
# Compute descriptive stats
desc_stats <- data.frame(
  Min = apply(df.clust, 2, min),    # minimum
  Med = apply(df.clust, 2, median), # median
  Mean = apply(df.clust, 2, mean),  # mean
  SD = apply(df.clust, 2, sd),      # Standard deviation
  Max = apply(df.clust, 2, max)     # Maximum
)
desc_stats <- round(desc_stats, 1)
head(desc_stats)
```

	Min <dbl>	Med <dbl>	Mean <dbl>	SD <dbl>	Max <dbl>
administrative	0	0	0.1	0.1	1.0
administrative_duration	0	0	0.0	0.1	1.0
informational	0	0	0.5	1.3	24.0
informational_duration	-1	0	34.8	141.5	2549.4
productrelated	0	0	0.0	0.1	1.0
productrelated_duration	0	0	0.0	0.0	1.0
6 rows					

Hide

```
# scale the datafarm to make the variables comparable
# entails transforming variables such that they have a mean of 0 and std dev of 1
# because we dont want the hierachical clustering result to depend on an arbitrary
variable unit

df.clust <- scale(df.clust)
head(df.clust)
```

```
      administrative administrative_duration informational informational_duration p
roductrelated
[1,]      -0.7025315          -0.4601081      -0.3988128          -0.2462725
-0.6963635
[2,]      -0.7025315          -0.4601081      -0.3988128          -0.2462725
-0.6739424
[3,]      -0.7025315          -0.4657410      -0.3988128          -0.2533417
-0.6963635
[4,]      -0.7025315          -0.4601081      -0.3988128          -0.2462725
-0.6739424
[5,]      -0.7025315          -0.4601081      -0.3988128          -0.2462725
-0.4945739
[6,]      -0.7025315          -0.4601081      -0.3988128          -0.2462725
-0.2927843
      productrelated_duration  bouncerrates  exitrates  pagevalues  specialday operati
ngsystems      browser
[1,]          -0.6289343  3.954699721  3.4273070 -0.3190356 -0.3103105      -
1.2396607 -0.7939682
[2,]          -0.5955997 -0.450343788  1.2650121 -0.3190356 -0.3103105      -
0.1371074 -0.2093703
[3,]          -0.6294551  3.954699721  3.4273070 -0.3190356 -0.3103105
2.0679992 -0.7939682
[4,]          -0.6275453  0.650917089  2.1299300 -0.3190356 -0.3103105
0.9654459 -0.2093703
[5,]          -0.3020990 -0.009839437  0.1838646 -0.3190356 -0.3103105
0.9654459  0.3752276
[6,]          -0.5486101 -0.102577188 -0.3661929 -0.3190356 -0.3103105      -
0.1371074 -0.2093703
      region traffictype
[1,] -0.8962939 -0.76562243
[2,] -0.8962939 -0.51660683
[3,]  2.4336556 -0.26759123
[4,] -0.4800502 -0.01857564
[5,] -0.8962939 -0.01857564
[6,] -0.8962939 -0.26759123
```

[Hide](#)




```
# clustering
# use the R function hclust() for hierarchical clustering
# 1st use dist() function to compute the Euclidean Distance btwn observs
# d will be the 1st argument in the hclust() function dissimilarity matrix

d<-dist(df.clust,method = "euclidean")
res.hc <- hclust (d,method = "ward.D2")
```

[Hide](#)

```
# plot obtained dendrogram
plot(res.hc,cex = 0.6,hang = -1)
```

[Hide](#)

```
# checking for accuracy
mean(df.clust== result$cluster)
```

```
[1] 0
```